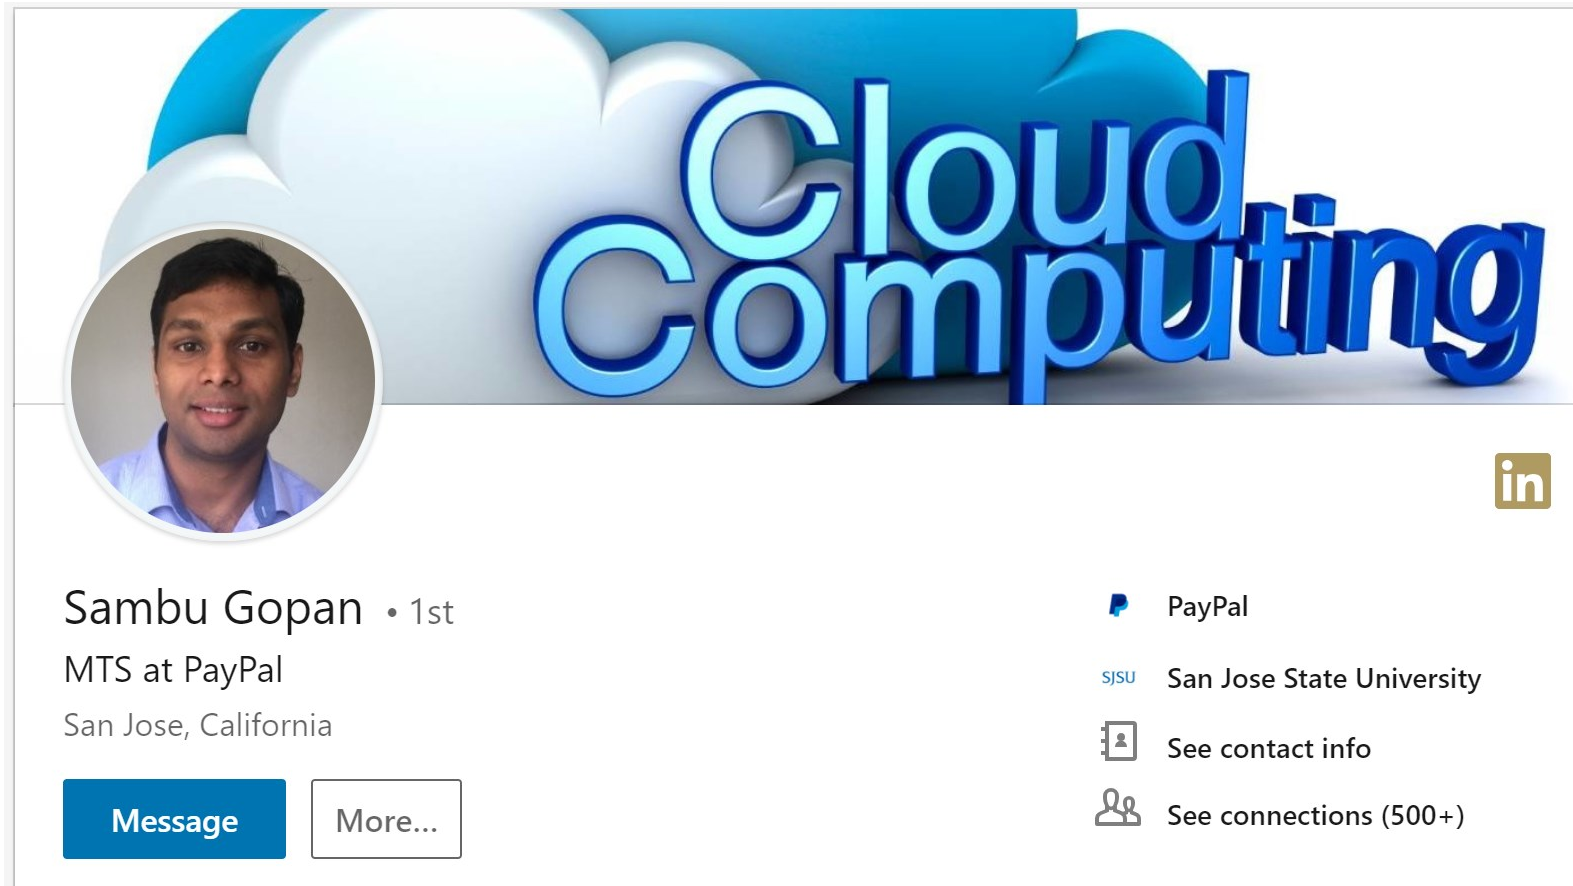



Guest Lecture 10/23





The image shows a LinkedIn profile for Sambu Gopan. The header features a large blue and white cloud graphic with the text "Cloud Computing" in a 3D blue font. To the left of the header is a circular profile picture of Sambu Gopan, a man with dark hair wearing a light blue shirt. Below the profile picture, the name "Sambu Gopan" is displayed with "• 1st" next to it, followed by "MTS at PayPal" and "San Jose, California". There are two buttons: a blue "Message" button and a white "More..." button with a black border. To the right of the profile information, there is a small brown "in" logo. Below the logo, there are four items: a PayPal logo with the text "PayPal", a San Jose State University logo with the text "San Jose State University", a contact icon with the text "See contact info", and a connections icon with the text "See connections (500+)".


Sambu Gopan • 1st
MTS at PayPal
San Jose, California

[Message](#) [More...](#)

 PayPal

 San Jose State University

 See contact info

 See connections (500+)

Introduction to MongoDB (NoSQL DB)

Why NoSql ?

- Relational databases are not designed to scale
- schema, joins

C and Latency Tradeoff

- Amazon claims that just an extra one tenth of a second on their response times will cost them 1% in sales.
- Google said they noticed that just a half a second increase in latency caused traffic to drop by a fifth.

4 Key Words on NoSQL

- Scale
- Speed
- Cloud
- New Data

What is NoSQL?

- non-relational
- simple API
- schema-free
- open-source
- horizontally scalable (sharding)
- replication support
- eventually consistent /BASE

Different types of NoSQL Databases

- NoSQL database are classified according to their data storage models:
 - Column (Cassandra)
 - Document (MongoDB)
 - Key – value Pair(Dynamo – Amazon)
 - Graph

MongoDB

- Name derived from Hu(**MONGO**)us word
- Document Oriented Database
- Built for High – Performance and scalability
- Document based queries for **Easy Readability**
- Replication and failover for **High Availability**
- Auto Sharding for **Easy Scalability**

Comparison between RDBMS and NoSQL DB

- Example: Class
- Location
- Presenter
 - Presenting at a location
- People
 - Potential attendees in context of a class
- Class
 - Presenter in location with people as actual attendees

Relational Database: Example

- Class schema in a relational database
- Presentation { id, name, location }
- People { id, name }
- Address { id, city, state, zip }

Schema for this class in a relational database model

Presentation			Address		
id name location			id city state		
1	Chris	SJSU	SJSU	San Jose	CA

People		Class	
id name		id person presentation	
10	Simon	20	10
11	Chris	20	11

Relational database: Example

```
CREATE TABLE Presentation (  
    id Integer primary key, name String, location string,  
    FOREIGN KEY (location) REFERENCES Address(id));  
CREATE TABLE Address (  
    id String primary key, city String, state String);  
CREATE TABLE People (  
    id Integer primary key, name String);  
CREATE TABLE Class (  
    id Integer, person Integer, presentation Integer,  
    PRIMARY KEY (id, person, presentation),  
    FOREIGN KEY (person) REFERENCES People(id),  
    FOREIGN KEY (presentation) REFERENCES Presentation(id));
```

Relational database: Example

```
select Presentation.name, Presentation.location,  
       Address.city, Address.state, People.name  
from Presentation, Address, People, Class  
where Class.person = People.id  
       and Class.presentation = Presentation.id  
       and Presentation.location = Address.id;
```

name	location	city	state	name
Chris	SJSU	San Jose	CA	Simon
Chris	SJSU	San Jose	CA	Chris

Relational Database: Recap

1. Schema design

Primary key (underlined) and foreign key (*cursive*) constraints

2. Table creation

DDL

3. Data insertion for each table

DML

4. Query: join

DML

5. Data structure creation within application system

JDBC resultset to e.g. Java objects

NoSQL Database: Use Case Example

```
use course /* database will be created if not present */
db.presentation.insert(
  {"id": 1,
    "name": "Simon",
    "location": {"id": "SJSU",
                  "city": "San Jose",
                  "state": "CA"},
    },
  "people": [{"id": 10, "name": "Simon"},
              {"id": 11, "name": "Chris"}
            ]
  })
```

NoSQL Database: Use Case Example

- `db.presentation.find()`
- `db.presentation.find({"id": 1})`

NoSQL Database: Recap

1. Schema design

Primary key (underlined) and foreign key (cursive) constraints

2. Table creation

DDL

3. Data insertion for each table

DML

4. Query: join

DML



5. Data structure creation within application system

IDBC resultset to e.g. Java objects

NoSQL Database: Major Players

- Too many document NoSQL databases to name a few distinct ones

29 systems in ranking, July 2014

Rank	Last Month	DBMS	Database Model	Score	Changes
1.	1.	MongoDB	Document store	238.78	+7.33
2.	2.	CouchDB	Document store	23.07	+0.28
3.	3.	Couchbase	Document store	16.58	+0.79
4.	4.	MarkLogic	Multi-model 	8.20	-0.02
5.	5.	RavenDB	Document store	5.09	-0.42
6.	6.	GemFire	Document store	2.16	-0.06
7.	7.	OrientDB	Multi-model 	1.71	-0.02
8.	8.	Cloudant	Document store	1.70	+0.07
9.	9.	Datameer	Document store	0.88	+0.08
10.	10.	Mnesia	Document store	0.72	+0.01

Key Benefit of NoSQL: $O(1)$ Lookup

- Fast lookup
 - No joining required
 - All data about one domain concept in one document
- Direct programming language representation
 - No mapping or ‘ORM’ layer required
- JSON library
 - Direct result representation and manipulation
 - JavaScript: representation in language data types directly
 - E.g., check out MongoDB node.js driver

Key Problem of NoSQL: No Join Operator

- Many NoSQL databases do not implement a join query operator
 - If you need to join data, then you have to do it in the application system layer
- But, wait a moment ...
 - Is it ever necessary to join data in NoSQL databases?
 - Some claim: not necessary due to support of
 - Sub-documents
 - Arrays (lists)
- Let's look at an example
 - Supplier - Parts

Key Problem of NoSQL: No Join Operator

- Example
 - Supplier - Parts relationship (N:M)
 - Each supplier supplies many parts
 - Each part supplied by many suppliers
- Relational DBMS
 - “Supplier” table
 - “Part” table
 - “Supplies” relationship in table

Key Problem of NoSQL: No Join Operator

Supplier - Part - Supplies

Supplier		Part		Supplies	
id	name	id	name	<i>supplier_id</i>	<i>part_id</i>
10	Supp1	20	Part1	10	20
11	Supp2	21	Part2	10	21
				11	20

Key Problem of NoSQL: No Join Operator

Supplier - Supplies – Part

```
{ "id": 10,  
  "name": "Supp1",  
  "supplies": [{ "id": 20, "name": "Part1"},  
                { "id": 21, "name":  
                  "Part2"} ] }  
  
{ "id": 11,  
  "name": "Supp2",  
  "supplies": [{ "id": 20, "name": "Part1"} ] }
```

Supplier - Supplies – Part

```
{ "id": 10,  
  "name": "Supp1",  
  "supplies": [20, 21] }  
  
{ "id": 10,  
  "name": "Supp1",  
  "supplies": [20, 21] }  
  
{ "id": 20, "name": "Part1" }  
{ "id": 21, "name": "Part2" }
```

Why use MongoDB?

- MongoDB stores data in Objects
- Uses BSON (Binary JSON)
- No Joins
- No Complex Queries
- Embedded Documents and arrays reduce the need for joins
- No multi-document transactions

Where to use MongoDB ?

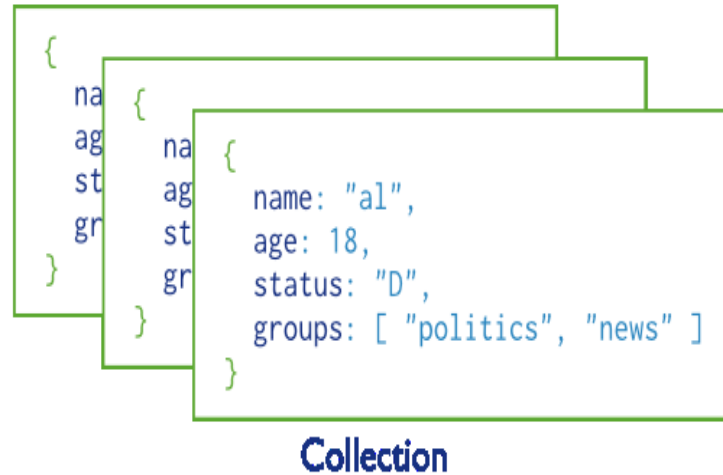
- Ideal for Web Applications
- Applications containing semi-structured data and need flexible schema management
- Caching and High Scalability
- Scenarios where **data availability** and **size of data** are priorities over the **transactions** of data

Terminology

- Mysql
- Table
- Row
- Column
- Joins
- Group By
- MongoDB
- Collection
- Document
- Field
- Not Recommended (\$lookup)
- Aggregation

Collections in MongoDB

- MongoDB stores all data in Collections
- It is schema – less and contains a group of related documents
- Created on-the-fly when referenced for the first time



Document in MongoDB

- Stored in Collections
- Has **_id** field – works like Primary keys in Relational databases
- Sample document containing name, age, status and groups

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

Queries in MongoDB

- MongoDB provides **db.collection.find()** method
- This method accepts both query criteria and projections

- `db.users.find(`
 `{ age: { $gt: 18 } },`
 `{ name: 1, address: 1 }`
 `).limit(5)`
 - ← **collection**
 - ← **query criteria**
 - ← **projection**
 - ← **cursor modifier**

Projections - Queries in MongoDB

- If you include 1 –it returns the value
- If you include 0 –it eliminates it from the result

```
db.records.find( { "user_id": { $lt: 42} }, { "_id": 0, "name": 1 , "email": 1 } )
```

- `_id` – always included in results. Specify “`_id : 0`” to exclude it from results

Insert Operation

- In MongoDB, `db.collection.insert()` method adds new documents to collections

```
db.users.insert (  ← collection
{
  name: "sue",      ← field:value
  age: 26,          ← field:value
  status: "A"       ← field:value
}                  } document
)
```

Update Operation

- In MongoDB, `db.collection.update()` method modifies existing documents in a collection

```
db.users.update(  
  { age: { $gt: 18 } },  
  { $set: { status: "A" } },  
  { multi: true }  
)
```

← collection
← update criteria
← update action
← update option

Remove Operation

- In MongoDB, `db.collection.remove()` method deletes document from the collection

```
db.users.remove(  
  { status: "D" }  
)
```

← collection

← remove criteria

Using MongoDB with Node.js


- **Install MongoDB Node.js Module**

```
npm install mongodb
```

Example

- Login Application
- Access MongoDB to authentic the user
- Use Mongo Store to store sessions in
MongoDB

Example – Login Page



Username:

Password:

Login

Exercise

- Create a Login Application
- Should have option to sign up the user
- Login with the same user
- Show the details of the signed in user
- Use MongoDB to store the data
- Use Passport with Express session for authentication

References

- SQL vs NoSQL - <https://www.mongodb.com/nosql-explained>
- MongoDB Introduction - <http://docs.mongodb.org/manual/core/crud-introduction/>
- Installing MongoDB (Mac) - <https://www.youtube.com/watch?v=WJ8m5QHvwc>
- Installing MongoDB (Windows) - <https://www.youtube.com/watch?t=1&v=sBdaRlgb4N8>

When to not use MongoDB?

- ACID properties are important for storage
- Highly Transactional Applications (Banking domain, Security)
- Problems and applications requiring Joins and complex queries

Key Problem of NoSQL: No Database-Enforced Consistency

- Not enforced
 - Primary key
 - Foreign key
 - Enumeration
 - Cascading delete
 - etc.
- Enforcement can be accomplished
 - When
 - reading or writing
 - In application system code
 - In self-implemented database access layer
 - In separate consistency check process
 - Not at all

How does MongoDB Store data?

- Stores data in form of Documents
- JSON like field – value pair
- Documents analogous to structures in programming languages with key – value pair
- Documents stored in **BSON (Binary JSON)** format
- BSON is JSON with additional type information

NoSQL: Key Insights

- Specialized data models
 - Not universal, but optimized towards special cases
- Specialized query access
 - Not universal, but optimized towards special cases
- Different / absent consistency supervision
 - Relaxed constraints
- Trade-off
 - Gain through specialization
 - Implementation of missing functionality outside of database